DRAFT v0 (2026-10-22) - working title

Diagrams Don't Fly: Handles, Not Architecture Families, Decide Whether Small Foundation Models Can Plan for a Drone

Status: v0, structure and argument. Numbers marked [TBD] wait for the full D-166 matrix (three seeds) and the handle ablation (not built yet). Findings and their evidence are tracked in `docs/RESEARCH.md`. Nothing here is a claim until it appears there with evidence.

Abstract (draft)

Foundation-model autonomy for drones is usually described by its architecture family. The families are: - an LLM that calls skills; - a VLM that points to waypoints; - a vision-language-action (VLA) model that flies directly; - hierarchies that combine them.

We ask whether the family decides success. We build twelve such architectures as configurations of one runtime, with the same small on-board model (Qwen3-VL-4B, plus a trained 4B VLA adapter), the same real-time loop, the same path planner and safety shield, and the same twelve capability tasks. Built only from their diagram's modules, every model-driven architecture solves simple see-and-go missions and fails every task that needs planning: a conditional branch, an ordered visit, recovery from a closed passage, following a moving target (0/12 planning passes for every model-driven architecture over 432 flights).

The same model improves from 0/9 to 5/9 when only the connective logic changes. That logic includes typed task contracts, a single owner of mission sequencing, evidence-based completion, and the perception-to-state hand-off. We call this logic the system's *handles*. Adding three handles to the most complete vanilla stack, with no other change, raises it from 21/36 to 30/36 and from 0/12 to 6/12 planning passes. Removing the typed contract or the executive alone returns it to 21/36. The same handles also lift the two other planner architectures (C13 22 -> 27, C10 16 -> 21), and handles that fit the other connections lift every architecture without a planner (C2 12 -> 19 ... VLA alone 0 -> 6), though only planner stacks gain planning tasks. Two time-critical tasks stay unsolved by every architecture, the oracle included: they are bound by model latency, not by the logic between the modules.

We argue that the handles, not the architecture family, determine whether a small foundation model can plan for a drone. Architecture comparisons should report them as first-class factors.

1 Introduction (draft)

- Drone autonomy papers present systems as boxes: planner, perception, executor, monitor. Comparisons between families are made across different models, simulators, sensors and prompts, so a difference in success cannot be attributed to the family.
- We remove those confounds. One model, one runtime and one task suite run all families as configurations; only the architecture switches change.
- The first result is negative and, we think, important. The diagram alone does not plan: every vanilla configuration fails the planning tasks.
- The working systems we and others built (AerialClaw, our drone_control stack, and our own live stack after repair) all work because of logic outside the diagram. This matches the scaffold literature (LLM-Modulo, AgentSpec, Inner Monologue, ProgPrompt, SayCan, TypeFly).

- Contributions: 1. a controlled testbed where architecture families are configurations of one runtime and one model; 2. a twelve-architecture by twelve-task evaluation with paired seeds, an oracle ceiling and a cost vector; 3. a taxonomy of handles (H1-H8) with controlled add-one / remove-one ablations [TBD]; 4. evidence that a trained VLA is a good pilot inside a hierarchy and a poor planner alone.

2 Related work (outline)

- **Architecture families for LLM/VLM drones.** TypeFly, AerialClaw, See-Point-Fly, OnFly, PMR / Selective Agentic Recovery, CognitiveDrone (-R1), FLIGHT, AeroVLA. Source: `docs/research/*`, the 94-paper UAV survey.
- **Scaffolds and handles.** LLM-Modulo, AgentSpec, Embodied Agent Interface, Inner Monologue, ProgPrompt, Code as Policies, SayCan, CoPAL, the DroneServer MCP harness. Source: `docs/research/HANDLES_20261022.md`.
- **Benchmarks.** AerialVLN, UAV-Flow, CognitiveDrone. The gap: no architecture comparison under identical constraints (survey Table 8).

3 Testbed (draft, from D-159 - D-166)

- **One model.** Qwen3-VL-4B (NF4) serves every language / vision role. The D207 LoRA adapter (trained on UAV-Flow, 70.3 % on its test set) serves VLA pilot roles. One GPU owner; calls are serialized.
- **One runtime.** A wall-clock 20 Hz control loop. Model latency is paid, not simulated. Actions expire. All roles share the asynchronous scheduler.
- **Shared execution.**
- SUPER local path planning on the range fan, with the same parameters as the oracle ceiling;
- a velocity controller;
- a collision shield, except C7 (no shield by definition).
- **Tasks.** Twelve capability tasks in a rendered 3-D grid simulator with an independent evaluator that no model sees:
- basic: coordinate, visible target, approach-and-hover;
- perception: search, the overturned vehicle among two;
- planning: conditional gate, ordered visit, closing passage, follow;
- obstacles: post, slalom, low wall.
- **Protocol.** Development seeds 1060-1062 for every reported matrix. Held-out seeds run once. No per-task tuning, no retries.

4 Architectures as configurations (draft table)

ID	Who decides the goal	Pilot	Verifier	Monitor	Recovery	Memory
C0	evaluator goal (privileged ceiling)	classical	yes	-	-	-
C1	text-only LLM picks one skill per turn	classical	yes	-	-	skill history
C2	VLM points in the image	classical	-	-	-	-
C3	as C2	classical	yes	-	-	-
C4	as C2	classical	yes	periodic	-	-
C5	as C2	classical	yes	periodic	-	yes
C6	as C2	classical	yes	-	event	-
FULL	LLM plans steps, VLM grounds them	classical	yes	arrival + periodic	event	yes
C7	the VLA writes its own goal line	VLA	-	-	-	goal memory
C8	as C7	VLA + shield	yes	-	-	goal memory
C10	as FULL, planner on a period	VLA	yes	yes	periodic	yes
C13	as FULL, planner on events	VLA	yes	yes	event	yes

5 Results (to fill)

- **5.1 Vanilla matrix** (432 flights, seeds 1060-1062, zero collisions; D-169). Passes / 36: C0 28; C13 22; FULL 21; C5 and C6 18; C10 16; C3 and C4 15; C2 12; C1 9; C7 and C8 0. Planning tasks: 0/12 for every model-driven architecture (C0 4/12). Tables: `reports/paper_matrix_d165/tables.md`.
- **5.2 The same model, different handles.** The history D-143 -> D-145 (0/4 -> 4/4) and D-150 -> D-158 (0/9 -> 5/9).
- **5.3 Handle ablation** (FULL base, 3 seeds; D-171 v2, D-173 v3). v3: FULL 21/36; + H1 24/36;
- H2 21/36; + H1 and H2 27/36, with planning 0/12 -> 3/12 (ordered visit 3/3) and every non-planning task passed. The executive alone does nothing; the pair is better than either. Remaining failures: wall openings cannot be grounded by the 4B VLM, and tracking.
- **5.4 Stacking a perception handle** (D-174, D-175). + H5 measured openings: 28/36 when the measured facts are only given to the VLM, which ignores them; 30/36 when sensor-answerable questions are routed to the range sensor. The if-then gate goes to 3/3.
- **5.5 Remove-one ablation of the final stack** (D-180, 3 seeds x 12 tasks).

Arm	Passes / 36	Planning / 12
H1 + H2 + H5	30	6
minus H5	27	3
minus H2	21	0
minus H1	21	0
no handles (FULL)	21	0

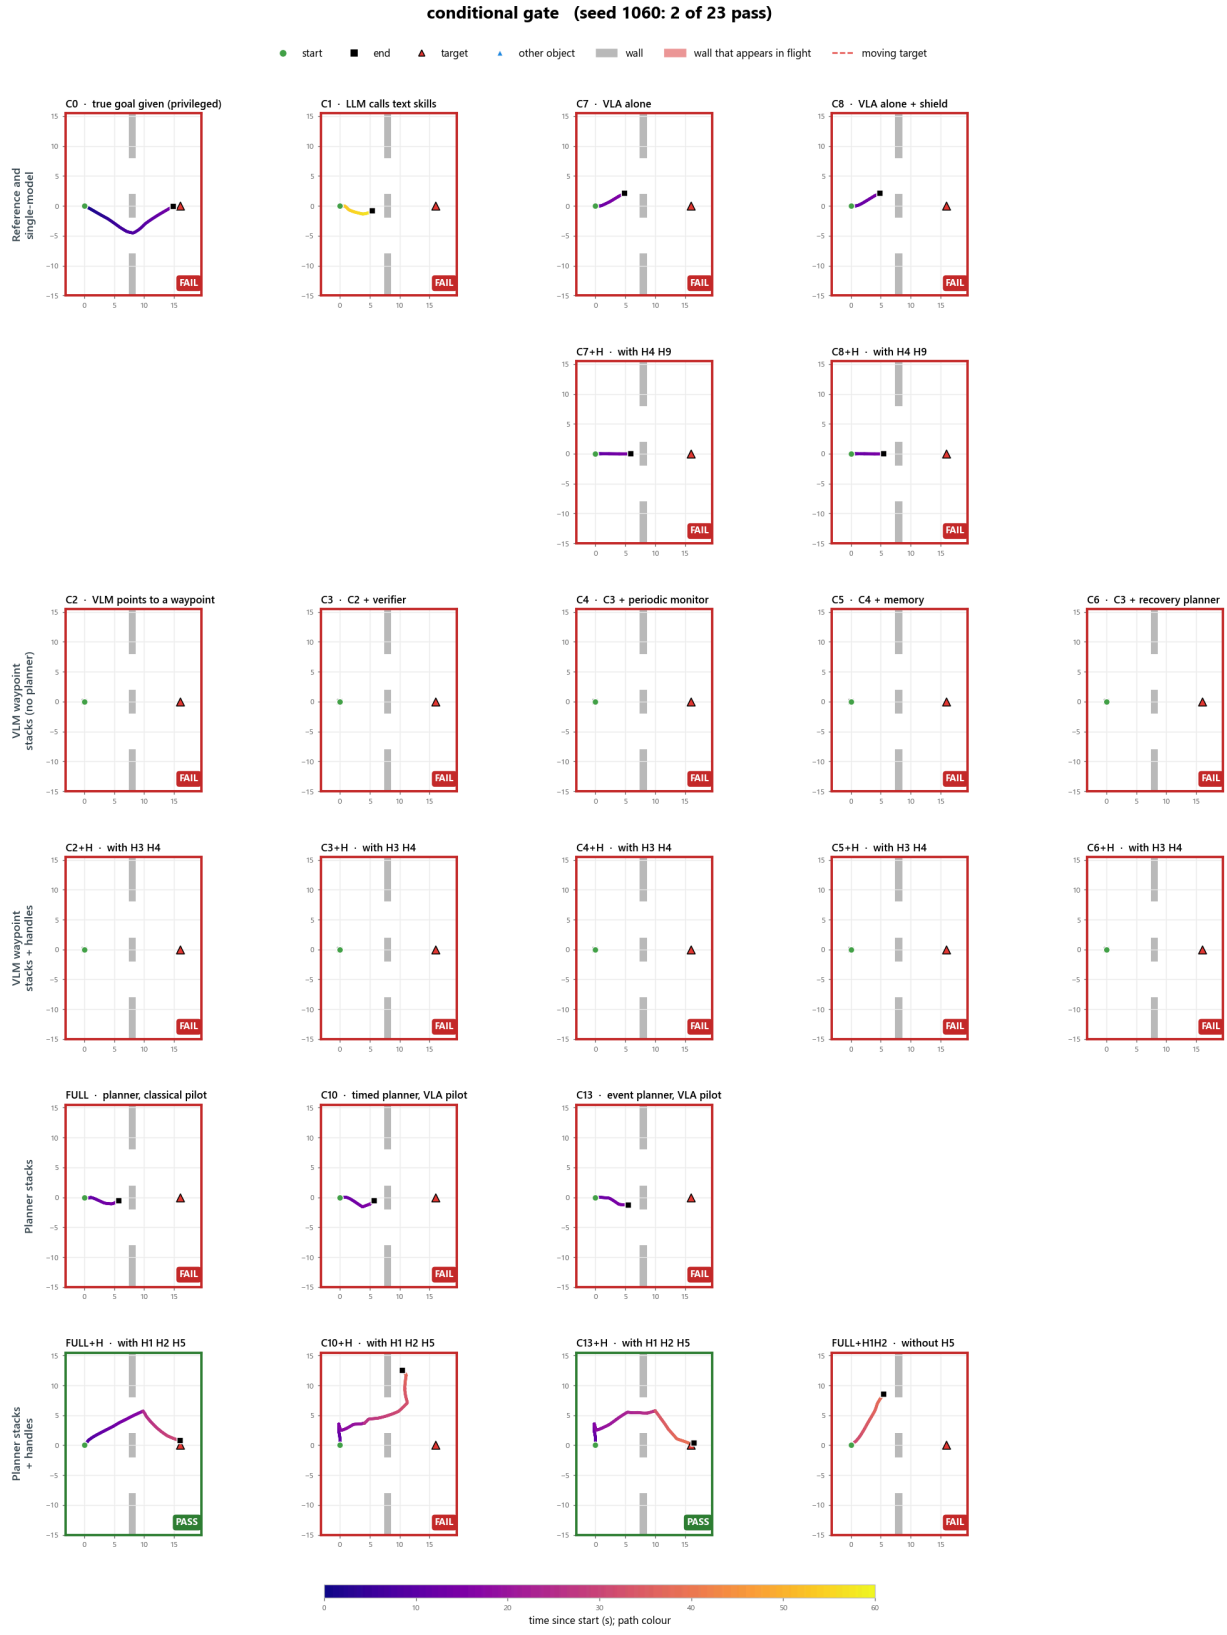
H1 and H2 are each necessary: removing either returns the stack to the no-handle level. H5 only helps on top of both. - **5.6 Negatives and the latency floor** (D-176 to D-179). Looking while moving (H7) and two wall-detour designs did not help. follow_target and closing_passage fail for every architecture, C0 included on follow, because the first move comes after about 5.7 s of model time. - **5.7 VLA alone vs in a hierarchy.** C13 22/36 against C7 / C8 0/36 (D-169). - **5.8 The same handles on other architectures** (D-181). The handle set fits unchanged on every architecture with a planner -> executor connection: FULL 21 -> 30, C13 22 -> 27, C10 16 -> 21. C13 + H solves the if-then gate as FULL + H does. On C10 the timed planner overrides the executive (two sequencing owners), and with the VLA pilot the ordered visit runs out of time. C1-C6 have no such

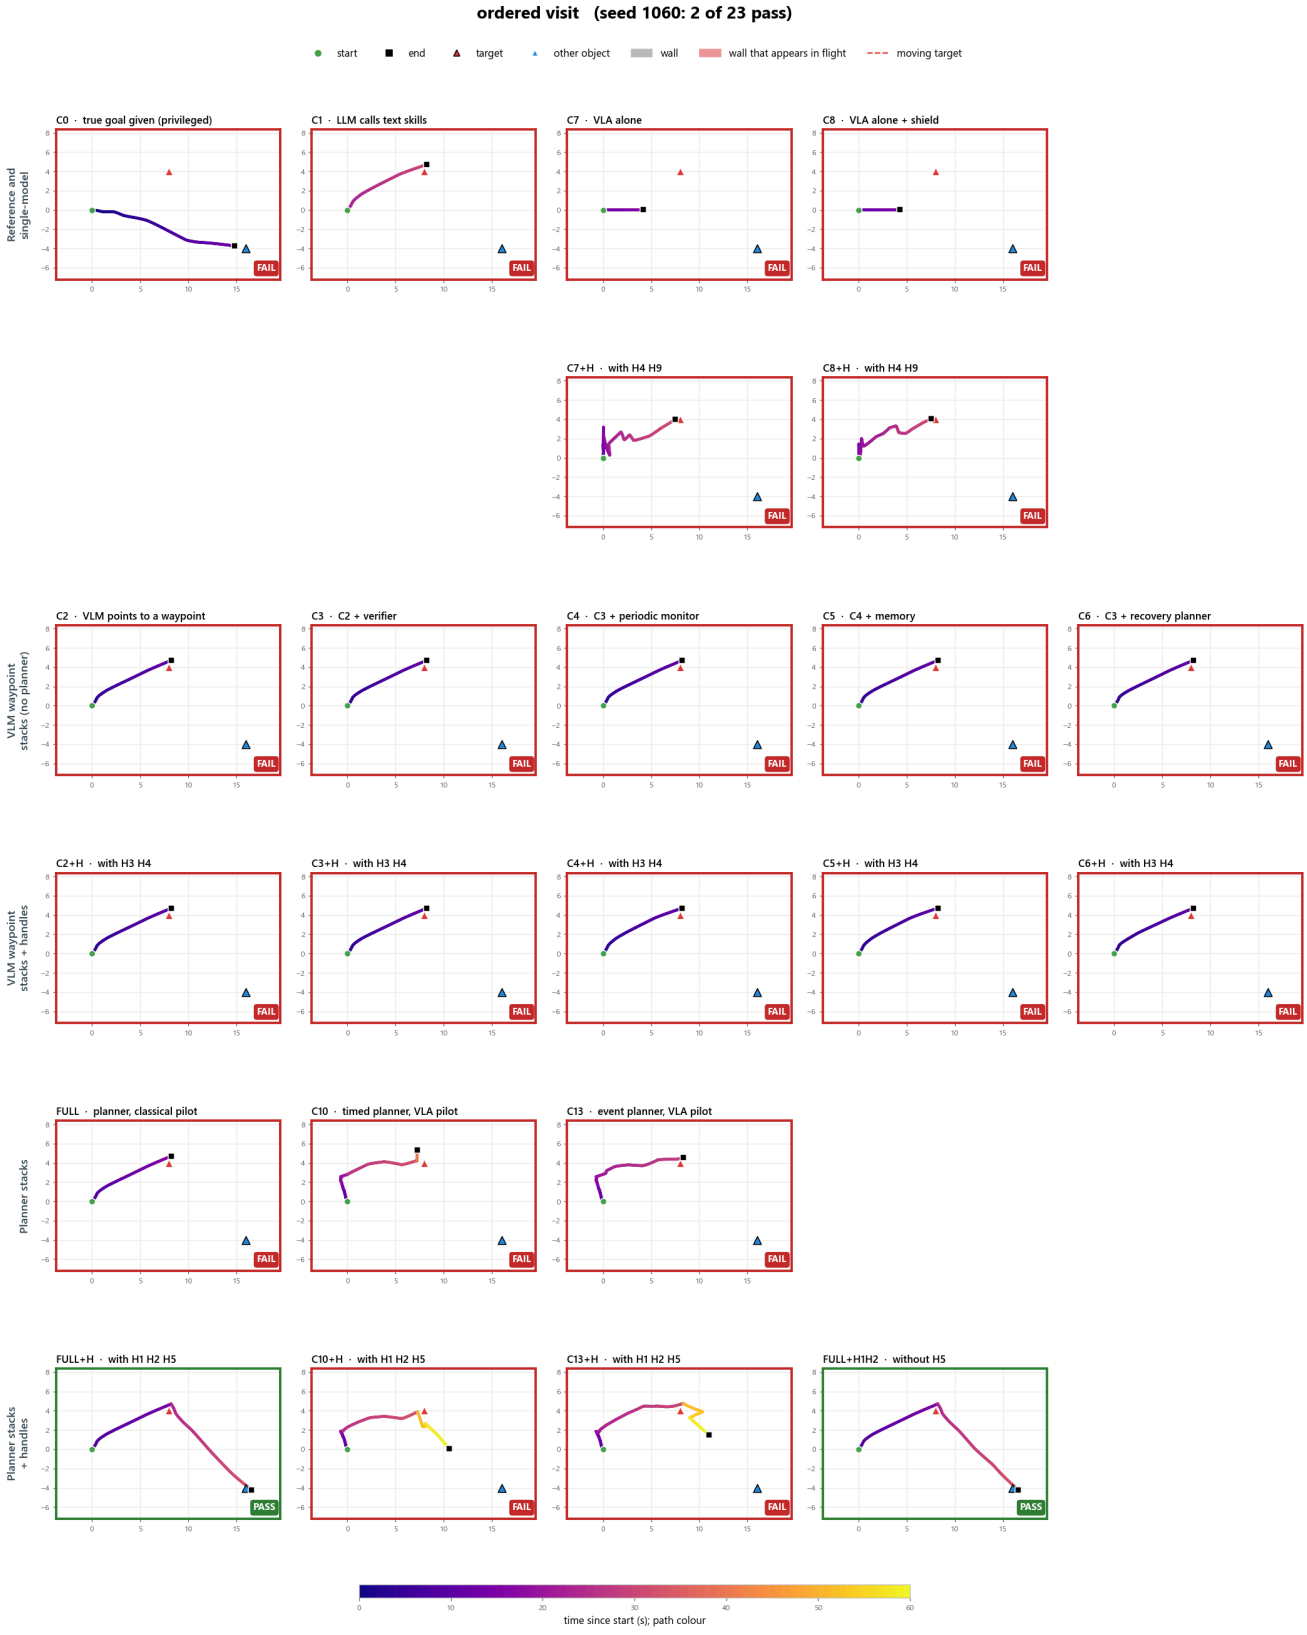
connection, so the handles are not applicable there; no box is added. - **5.9 Handles for the architectures without a planner** (D-182). H3 (search policy), H4 (evidence-based completion) and H9 (measured distance) act on connections C2-C8 already have. C2 12 -> 19, C3 15 -> 21, C4 15 -> 23, C5 18 -> 23, C6 18 -> 23, C7 0 -> 6, C8 0 -> 6. H3 and H4 already existed in FULL's planned path ("hidden handles") but never reached C2-C5. Planning tasks stay at 0 without a planner: the handles that solve planning need a plan connection to act on. - **Figures: task maps** (`scripts/paper_task_maps.py`). One page per task, one top-down panel per architecture, path coloured by time: - Figure M0: architecture legend (`figures/maps_legend.png`). - Figure M1: if-then gate (`figures/map_s1060_conditional_gate.png`). - Figure M2: ordered visit (`figures/map_s1060_ordered_visit.png`). - Figure M3: closing passage (`figures/map_s1060_closing_passage.png`). - Figure M4: follow target (`figures/map_s1060_follow_target.png`). - Figure M5: obstacle slalom (`figures/map_s1060_obstacle_slalom.png`). - Figure M6: approach and hover (`figures/map_s1060_approach_hover.png`). - Figure M7: turn and search (`figures/map_s1060_turn_search.png`).

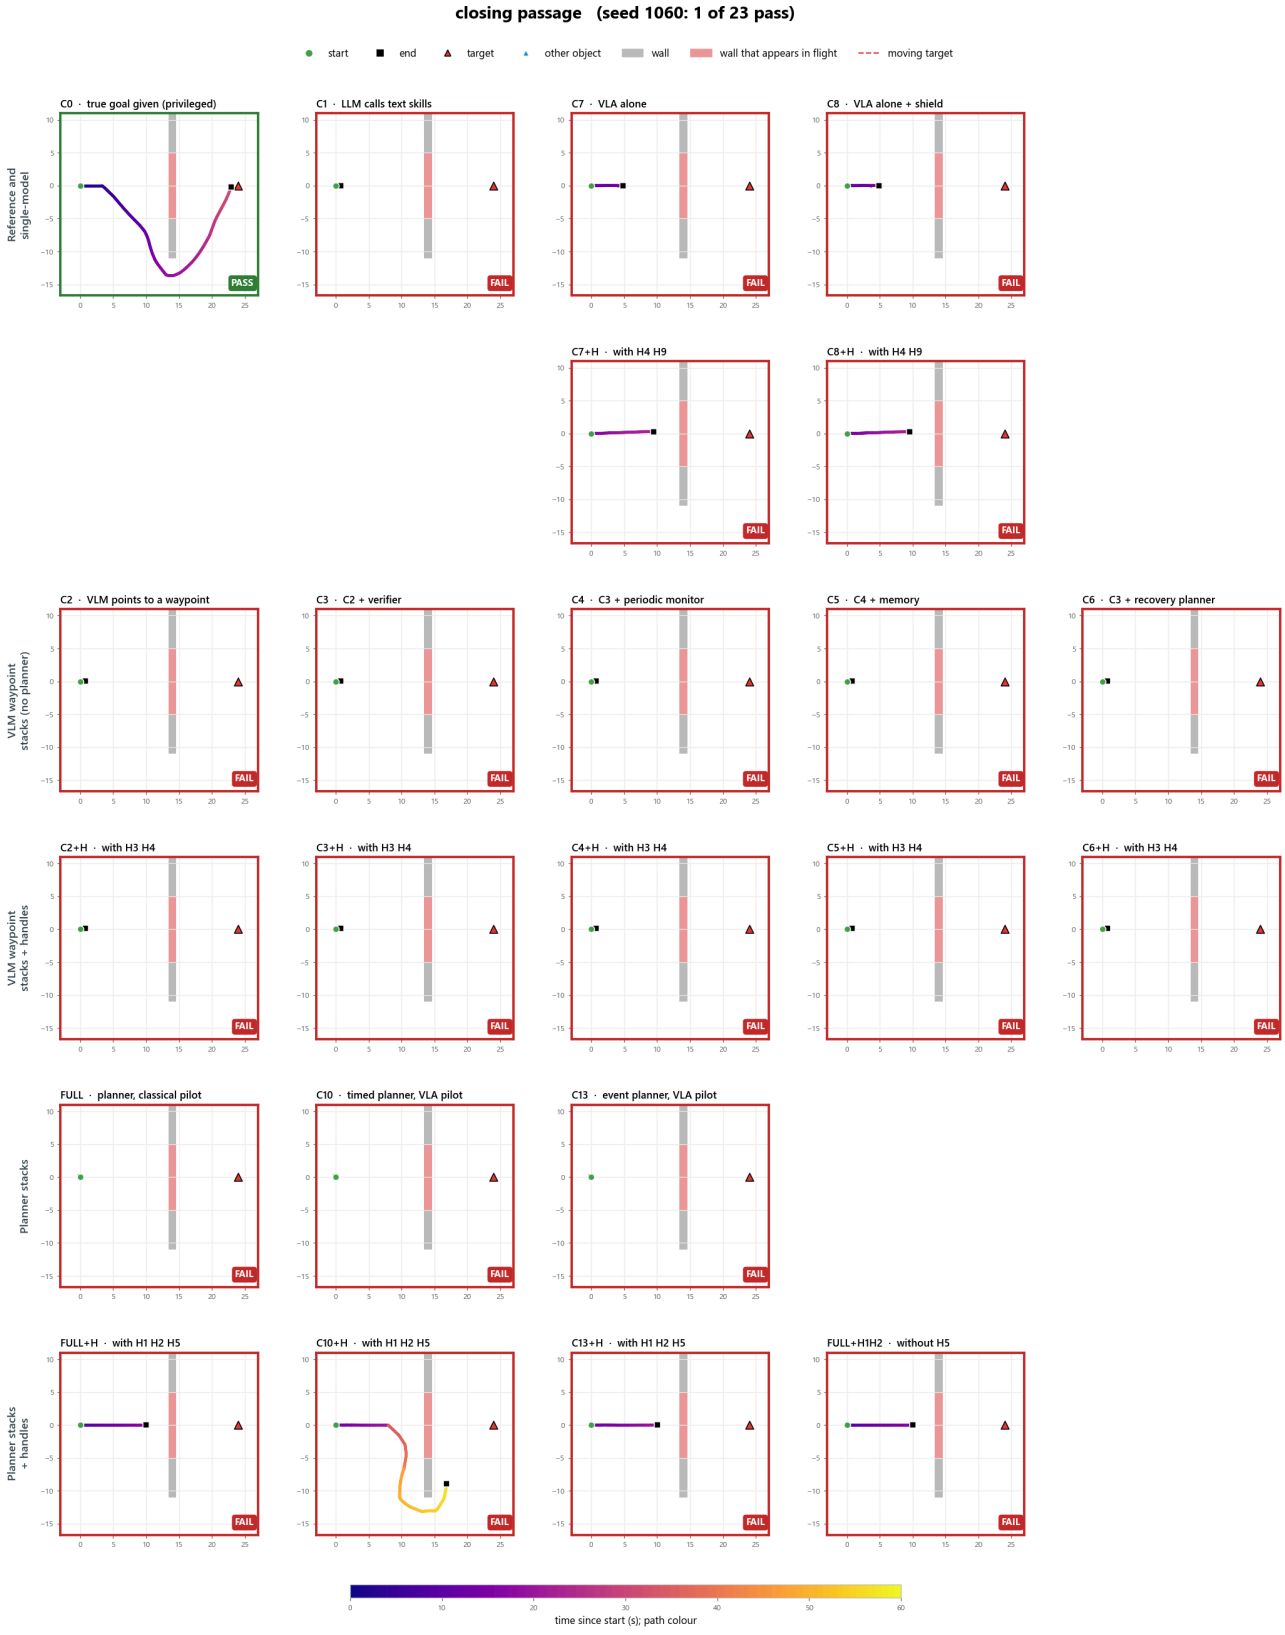
Architectures and handles in the task maps

C0	Oracle: the true goal is given to a classical pilot. Privileged ceiling.
C1	Skill agent: a text-only LLM calls goto / approach / detect / scan / hover.
C2	VLM points at the target in the image; the point becomes a waypoint.
C3	C2 plus a verifier that checks each waypoint before it is flown.
C4	C3 plus a monitor that reviews progress on a timer.
C5	C4 plus memory of past goals and outcomes.
C6	C3 plus a recovery planner called on events (lost target, stall).
C7	VLA alone: the trained VLA writes its own goal line and flies it.
C8	C7 plus the safety shield and envelope verifier.
FULL	Planner + VLM grounding + verifier + monitor + event recovery + memory; classical pilot.
C10	As FULL, the planner runs on a timer; the VLA is the pilot.
C13	As FULL, the planner runs on events; the VLA is the pilot.
H1	Handle: typed task contract (visit / pass / condition / follow / goto / search).
H2	Handle: one executive runs the contract by step kind.
H5	Handle: measured openings from the range sensor; sensor answers what it can.
H3	Handle: a find / search mission starts the bounded camera search itself.
H4	Handle: completion on evidence: stated standoff / hold, object passed below.

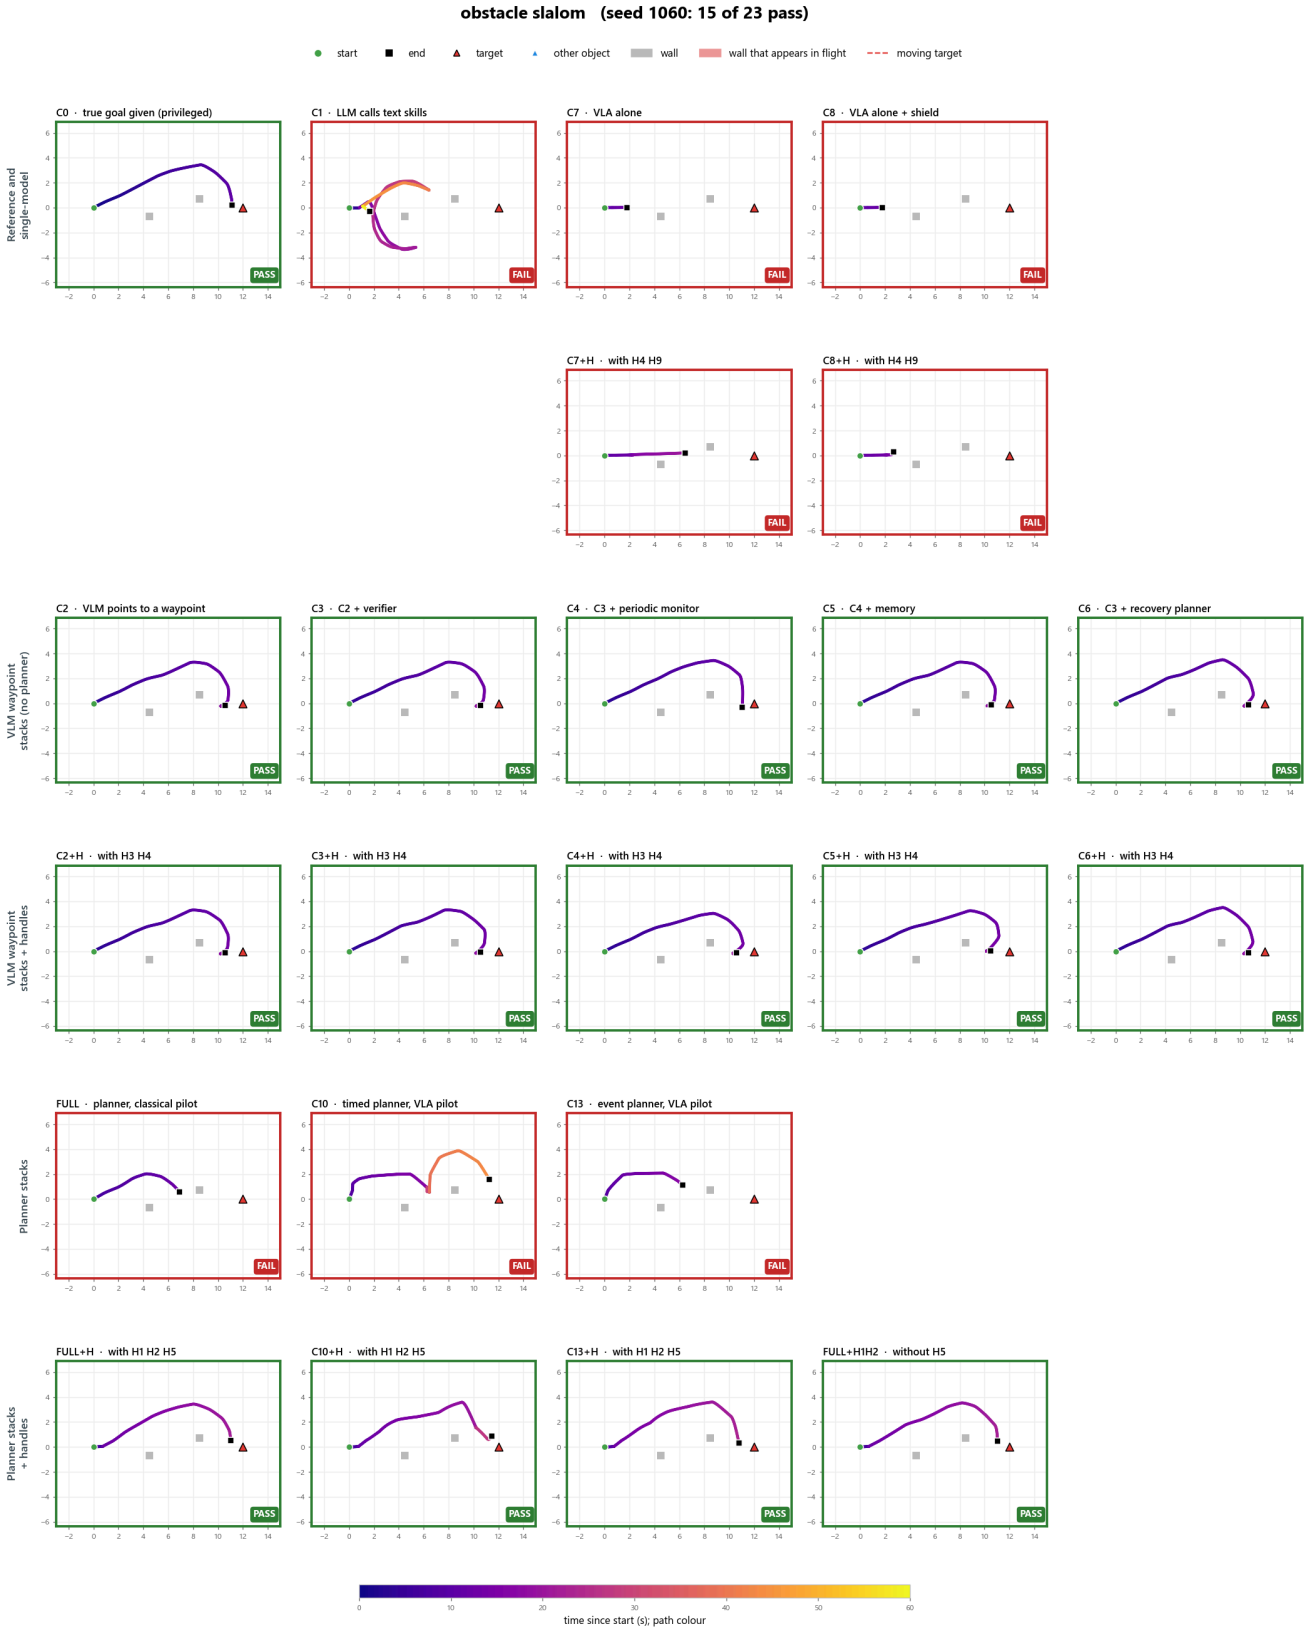
Same model (Qwen3-VL-4B + D207 VLA adapter), runtime, path planner, shield and tasks in every panel. '+H' = the handles that fit that architecture.


conditional gate (seed 1060: 2 of 23 pass)
start
end
target
other object
wall
wall that appears in flight
moving target
Reference and single-model
C0 · true goal given (privileged)
C1 · LLM calls text skills
C7 · VLA alone
C8 · VLA alone + shield
C7+H · with H4 H9
C8+H · with H4 H9
VLM waypoint stacks (no planner)
C2 · VLM points to a waypoint
C3 · C2 + verifier
C4 · C3 + periodic monitor
C5 · C4 + memory
C6 · C3 + recovery planner
VLM waypoint stacks + handles
C2+H · with H3 H4
C3+H · with H3 H4
C4+H · with H3 H4
C5+H · with H3 H4
C6+H · with H3 H4
Planner stacks
FULL · planner, classical pilot
C10 · timed planner, VLA pilot
C13 · event planner, VLA pilot
Planner stacks + handles
FULL+H · with H1 H2 H5
C10+H · with H1 H2 H5
C13+H · with H1 H2 H5
FULL+H1H2 · without H5
FAIL
PASS
time since start (s); path colour

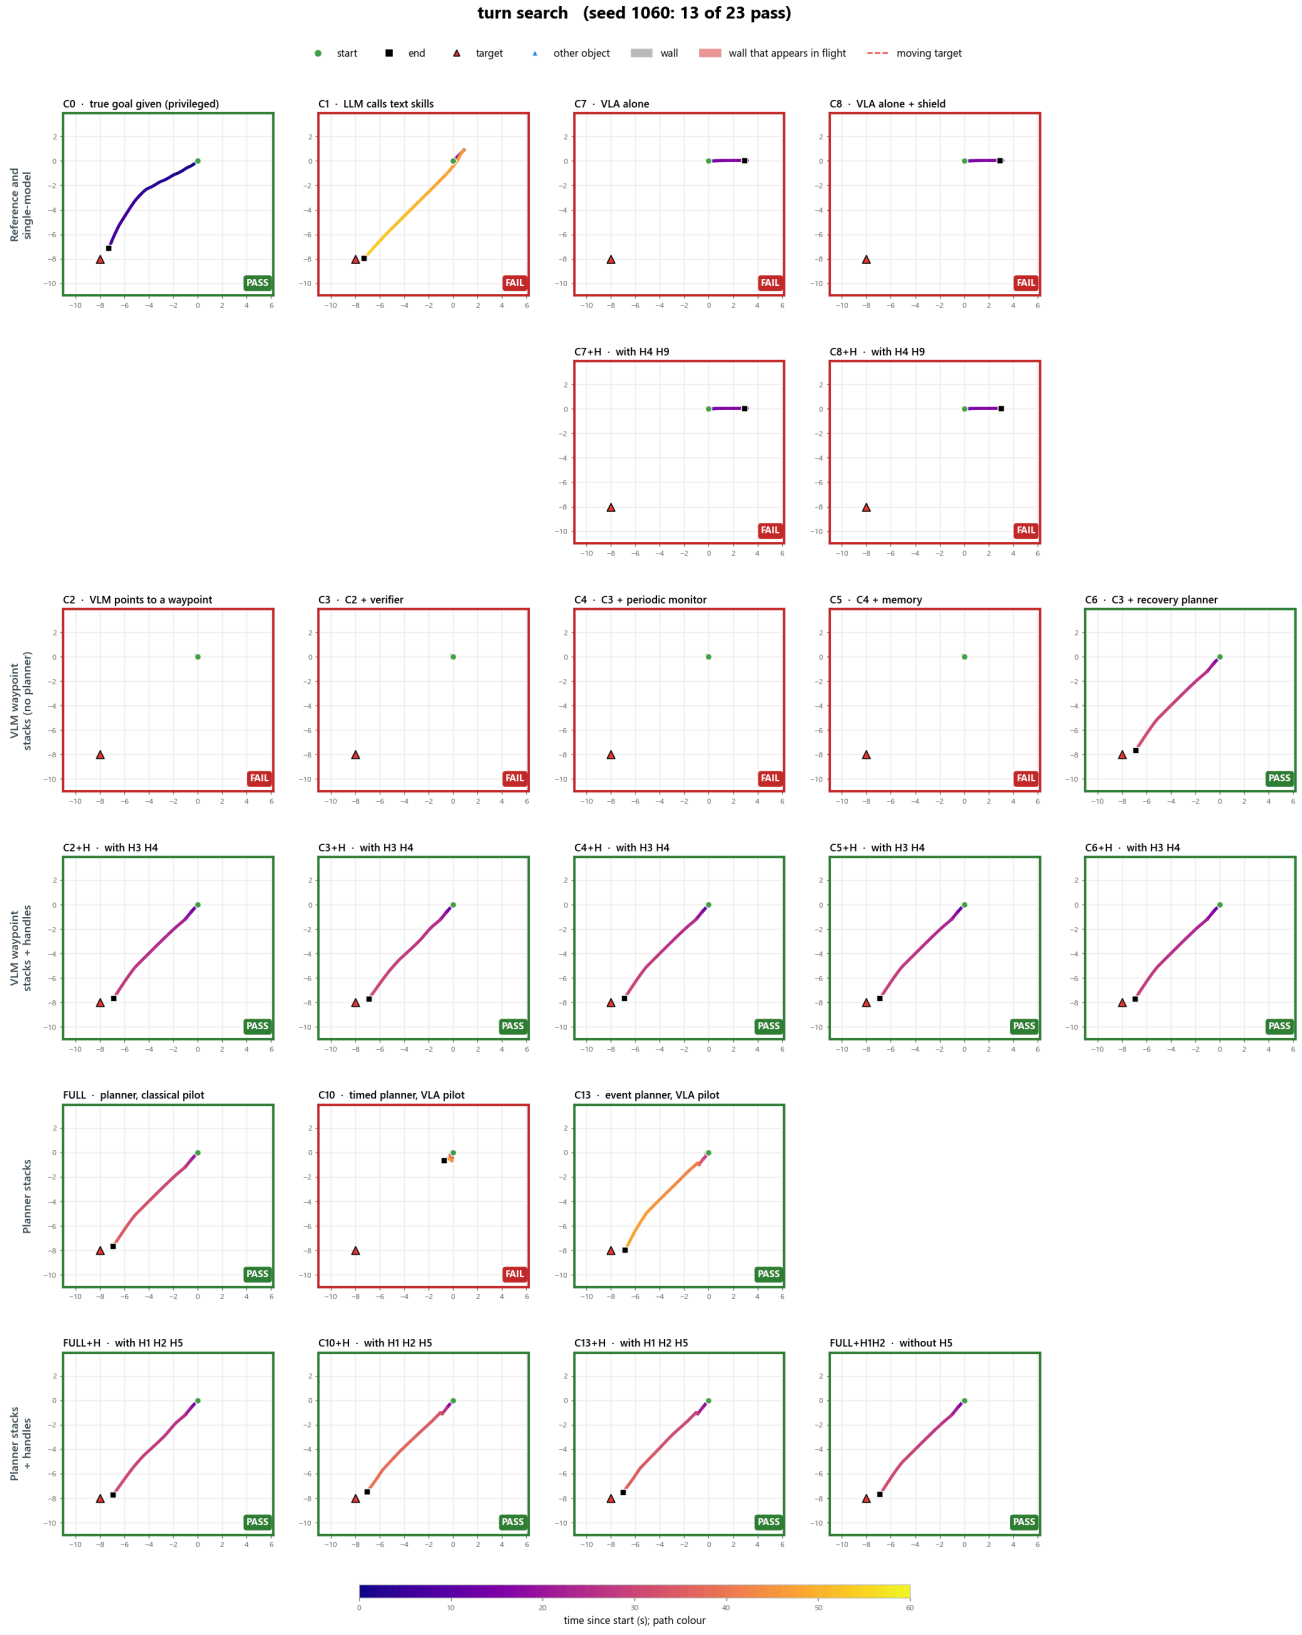

ordered visit (seed 1060: 2 of 23 pass)
start
end
target
other object
wall
wall that appears in flight
moving target
Reference and single-model
C0 · true goal given (privileged)
C1 · LLM calls text skills
C7 · VLA alone
C8 · VLA alone + shield
C7+H · with H4 H9
C8+H · with H4 H9
VLM waypoint stacks (no planner)
C2 · VLM points to a waypoint
C3 · C2 + verifier
C4 · C3 + periodic monitor
C5 · C4 + memory
C6 · C3 + recovery planner
VLM waypoint stacks + handles
C2+H · with H3 H4
C3+H · with H3 H4
C4+H · with H3 H4
C5+H · with H3 H4
C6+H · with H3 H4
Planner stacks
FULL · planner, classical pilot
C10 · timed planner, VLA pilot
C13 · event planner, VLA pilot
Planner stacks + handles
FULL+H · with H1 H2 H5
C10+H · with H1 H2 H5
C13+H · with H1 H2 H5
FULL+H1H2 · without H5
FAIL
PASS
time since start (s); path colour


closing passage (seed 1060: 1 of 23 pass)
start
end
target
other object
wall
wall that appears in flight
moving target
Reference and single-model
C0 · true goal given (privileged)
PASS
C1 · LLM calls text skills
C7 · VLA alone
C8 · VLA alone + shield
C7+H · with H4 H9
C8+H · with H4 H9
VLM waypoint stacks (no planner)
C2 · VLM points to a waypoint
C3 · C2 + verifier
C4 · C3 + periodic monitor
C5 · C4 + memory
C6 · C3 + recovery planner
VLM waypoint stacks + handles
C2+H · with H3 H4
C3+H · with H3 H4
C4+H · with H3 H4
C5+H · with H3 H4
C6+H · with H3 H4
Planner stacks
FULL · planner, classical pilot
C10 · timed planner, VLA pilot
C13 · event planner, VLA pilot
Planner stacks + handles
FULL+H · with H1 H2 H5
C10+H · with H1 H2 H5
C13+H · with H1 H2 H5
FULL+H1H2 · without H5
FAIL
time since start (s); path colour


follow target (seed 1060: 0 of 23 pass)
start
end
target
other object
wall
wall that appears in flight
moving target
Reference and single-model
C0 · true goal given (privileged)
C1 · LLM calls text skills
C7 · VLA alone
C8 · VLA alone + shield
C7+H · with H4 H9
C8+H · with H4 H9
VLM waypoint stacks (no planner)
C2 · VLM points to a waypoint
C3 · C2 + verifier
C4 · C3 + periodic monitor
C5 · C4 + memory
C6 · C3 + recovery planner
VLM waypoint stacks + handles
C2+H · with H3 H4
C3+H · with H3 H4
C4+H · with H3 H4
C5+H · with H3 H4
C6+H · with H3 H4
Planner stacks
FULL · planner, classical pilot
C10 · timed planner, VLA pilot
C13 · event planner, VLA pilot
Planner stacks + handles
FULL+H · with H1 H2 H5
C10+H · with H1 H2 H5
C13+H · with H1 H2 H5
FULL+H1H2 · without H5
FAIL
time since start (s); path colour


obstacle slalom (seed 1060: 15 of 23 pass)
start
end
target
other object
wall
wall that appears in flight
moving target
Reference and single-model
C0 · true goal given (privileged)
C1 · LLM calls text skills
C7 · VLA alone
C8 · VLA alone + shield
C7+H · with H4 H9
C8+H · with H4 H9
VLM waypoint stacks (no planner)
C2 · VLM points to a waypoint
C3 · C2 + verifier
C4 · C3 + periodic monitor
C5 · C4 + memory
C6 · C3 + recovery planner
VLM waypoint stacks + handles
C2+H · with H3 H4
C3+H · with H3 H4
C4+H · with H3 H4
C5+H · with H3 H4
C6+H · with H3 H4
Planner stacks
FULL · planner, classical pilot
C10 · timed planner, VLA pilot
C13 · event planner, VLA pilot
Planner stacks + handles
FULL+H · with H1 H2 H5
C10+H · with H1 H2 H5
C13+H · with H1 H2 H5
FULL+H1H2 · without H5
PASS
FAIL
time since start (s); path colour


approach hover (seed 1060: 14 of 23 pass)
start
end
target
other object
wall
wall that appears in flight
moving target
Reference and single-model
C0 · true goal given (privileged)
PASS
C1 · LLM calls text skills
FAIL
C7 · VLA alone
FAIL
C8 · VLA alone + shield
FAIL
C7+H · with H4 H9
FAIL
C8+H · with H4 H9
FAIL
VLM waypoint stacks (no planner)
C2 · VLM points to a waypoint
FAIL
C3 · C2 + verifier
FAIL
C4 · C3 + periodic monitor
FAIL
C5 · C4 + memory
PASS
C6 · C3 + recovery planner
FAIL
VLM waypoint stacks + handles
C2+H · with H3 H4
PASS
C3+H · with H3 H4
PASS
C4+H · with H3 H4
PASS
C5+H · with H3 H4
PASS
C6+H · with H3 H4
PASS
Planner stacks
FULL · planner, classical pilot
PASS
C10 · timed planner, VLA pilot
PASS
C13 · event planner, VLA pilot
PASS
Planner stacks + handles
FULL+H · with H1 H2 H5
PASS
C10+H · with H1 H2 H5
PASS
C13+H · with H1 H2 H5
PASS
FULL+H1H2 · without H5
PASS
time since start (s); path colour



6 Discussion (outline)

- **Why vanilla fails planning.** Per-module failure attribution, in the Embodied Agent Interface style:
- goal interpretation (the branch flattened in conditional_gate);
- state tracking (the second target lost in ordered_visit);
- recovery (no route when the passage closes);
- missing behaviour (no tracking skill).

- The taxonomy H1-H8 and what each fixes.
- Implications for comparing published systems: report the handles.

7 Limitations (outline)

- A synthetic rendered simulator, not AirSim or real flight.
- One model family, at 4B.
- Development seeds only, until the final held-out run.
- Task outcomes are near deterministic across the three development seeds (PPP or FFF), so each arm is closer to 12 outcomes than 36 trials, and the confidence intervals overstate precision.
- The handles are partly designed by us. We mitigate this by using only public mission text and sensed state, never evaluator truth.